

```
In [31]: import pandas as pd
d=pd.read_csv("/home/placement/Downloads/TelecomCustomerChurn.csv")
d.describe()
```

```
Out[31]:
```

	SeniorCitizen	tenure	MonthlyCharges
count	7043.000000	7043.000000	7043.000000
mean	0.162147	32.371149	64.761692
std	0.368612	24.559481	30.090047
min	0.000000	0.000000	18.250000
25%	0.000000	9.000000	35.500000
50%	0.000000	29.000000	70.350000
75%	0.000000	55.000000	89.850000
max	1.000000	72.000000	118.750000

```
In [32]: d.head()
```

```
Out[32]:
```

	customerID	gender	SeniorCitizen	Partner	Dependents	tenure	PhoneService	MultipleLines	InternetService	OnlineSecurity	...	DeviceProtection
0	7590-VHVEG	Female	0	Yes	No	1	No	No phone service	DSL	No	...	
1	5575-GNVDE	Male	0	No	No	34	Yes	No	DSL	Yes	...	
2	3668-QPYBK	Male	0	No	No	2	Yes	No	DSL	Yes	...	
3	7795-CFOCW	Male	0	No	No	45	No	No phone service	DSL	Yes	...	
4	9237-HQITU	Female	0	No	No	2	Yes	No	Fiber optic	No	...	

5 rows × 13 columns



```
In [33]: d.isna().sum()
```

```
Out[33]: customerID      0  
gender      0  
SeniorCitizen  0  
Partner      0  
Dependents    0  
tenure      0  
PhoneService  0  
MultipleLines  0  
InternetService  0  
OnlineSecurity  0  
OnlineBackup  0  
DeviceProtection  0  
TechSupport  0  
StreamingTV  0  
StreamingMovies  0  
Contract      0  
PaperlessBilling  0  
PaymentMethod  0  
MonthlyCharges  0  
TotalCharges  0  
Churn         0  
dtype: int64
```

In [34]: `d.info()`

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 7043 entries, 0 to 7042
Data columns (total 21 columns):
#   Column                Non-Null Count  Dtype  
---  -
0   customerID            7043 non-null   object 
1   gender                 7043 non-null   object 
2   SeniorCitizen          7043 non-null   int64  
3   Partner                7043 non-null   object 
4   Dependents             7043 non-null   object 
5   tenure                 7043 non-null   int64  
6   PhoneService           7043 non-null   object 
7   MultipleLines           7043 non-null   object 
8   InternetService        7043 non-null   object 
9   OnlineSecurity         7043 non-null   object 
10  OnlineBackup           7043 non-null   object 
11  DeviceProtection       7043 non-null   object 
12  TechSupport            7043 non-null   object 
13  StreamingTV            7043 non-null   object 
14  StreamingMovies        7043 non-null   object 
15  Contract               7043 non-null   object 
16  PaperlessBilling       7043 non-null   object 
17  PaymentMethod          7043 non-null   object 
18  MonthlyCharges         7043 non-null   float64 
19  TotalCharges           7043 non-null   object 
20  Churn                  7043 non-null   object 
dtypes: float64(1), int64(2), object(18)
memory usage: 1.1+ MB
```

In [35]: `#d1=d.drop(['customerID','gender','PaymentMethod','PaperlessBilling','Dependents','OnlineSecurity','OnlineBa
'TechSupport','StreamingTV','StreamingMovies','InternetService'],axis=1)`

In [36]: `d['Churn']=d['Churn'].map({'Yes':1,'No':0})`

```
In [37]: d['PhoneService']=d['PhoneService'].map({'Yes':1, 'No':0})
```

```
In [38]: d['Partner']=d['Partner'].map({'Yes':1, 'No':0})
```

```
In [39]: d['TotalCharges']=r.to_numeric(d['TotalCharges'], errors='coerce')
```

```
In [ ]:
```

```
In [40]: d.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 7043 entries, 0 to 7042
Data columns (total 21 columns):
#   Column                Non-Null Count  Dtype
---  -
0   customerID            7043 non-null   object
1   gender                7043 non-null   object
2   SeniorCitizen         7043 non-null   int64
3   Partner               7043 non-null   int64
4   Dependents            7043 non-null   object
5   tenure                7043 non-null   int64
6   PhoneService          7043 non-null   int64
7   MultipleLines         7043 non-null   object
8   InternetService       7043 non-null   object
9   OnlineSecurity        7043 non-null   object
10  OnlineBackup          7043 non-null   object
11  DeviceProtection      7043 non-null   object
12  TechSupport           7043 non-null   object
13  StreamingTV           7043 non-null   object
14  StreamingMovies       7043 non-null   object
15  Contract              7043 non-null   object
16  PaperlessBilling      7043 non-null   object
17  PaymentMethod         7043 non-null   object
18  MonthlyCharges        7043 non-null   float64
19  TotalCharges          7032 non-null   float64
20  Churn                 7043 non-null   int64
dtypes: float64(2), int64(5), object(14)
memory usage: 1.1+ MB
```

```
In [41]: #d=d.fillna(d.mean())
```

In [42]:

d

Out[42]:

	customerID	gender	SeniorCitizen	Partner	Dependents	tenure	PhoneService	MultipleLines	InternetService	OnlineSecurity	...	DevicePro
0	7590-VHVEG	Female	0	1	No	1	0	No phone service	DSL	No	...	
1	5575-GNVDE	Male	0	0	No	34	1	No	DSL	Yes	...	
2	3668-QPYBK	Male	0	0	No	2	1	No	DSL	Yes	...	
3	7795-CFOCW	Male	0	0	No	45	0	No phone service	DSL	Yes	...	
4	9237-HQITU	Female	0	0	No	2	1	No	Fiber optic	No	...	
...	
7038	6840-RESVB	Male	0	1	Yes	24	1	Yes	DSL	Yes	...	
7039	2234-XADUH	Female	0	1	Yes	72	1	Yes	Fiber optic	No	...	
7040	4801-JZAZL	Female	0	1	Yes	11	0	No phone service	DSL	Yes	...	
7041	8361-LTMKD	Male	1	1	No	4	1	Yes	Fiber optic	No	...	
7042	3186-AJIEK	Male	0	0	No	66	1	No	Fiber optic	Yes	...	

7043 rows × 21 columns



```
In [43]: d.isna().sum()
```

```
Out[43]: customerID      0
gender      0
SeniorCitizen  0
Partner      0
Dependents    0
tenure      0
PhoneService  0
MultipleLines  0
InternetService  0
OnlineSecurity  0
OnlineBackup  0
DeviceProtection  0
TechSupport    0
StreamingTV    0
StreamingMovies  0
Contract      0
PaperlessBilling  0
PaymentMethod  0
MonthlyCharges  0
TotalCharges   11
Churn          0
dtype: int64
```

```
In [44]: d=d.fillna(d.mean())
```

/tmp/ipykernel_13877/1862675393.py:1: FutureWarning: The default value of numeric_only in DataFrame.mean is deprecated. In a future version, it will default to False. In addition, specifying 'numeric_only=None' is deprecated. Select only valid columns or specify the value of numeric_only to silence this warning.

```
d=d.fillna(d.mean())
```

```
In [45]: y=d['Churn']
x=d.drop(['customerID', 'Churn'],axis=1)
```

```
In [46]: x.isna().sum()
```

```
Out[46]: gender          0  
SeniorCitizen          0  
Partner                0  
Dependents             0  
tenure                 0  
PhoneService           0  
MultipleLines          0  
InternetService        0  
OnlineSecurity         0  
OnlineBackup           0  
DeviceProtection       0  
TechSupport            0  
StreamingTV            0  
StreamingMovies        0  
Contract               0  
PaperlessBilling       0  
PaymentMethod          0  
MonthlyCharges         0  
TotalCharges           0  
dtype: int64
```



In [53]:

x

Out[53]:

	SeniorCitizen	Partner	tenure	PhoneService	MonthlyCharges	TotalCharges	gender_Female	gender_Male	Dependents_No	Dependents_Yes
0	0	1	1	0	29.85	29.85	1	0	1	(
1	0	0	34	1	56.95	1889.50	0	1	1	(
2	0	0	2	1	53.85	108.15	0	1	1	(
3	0	0	45	0	42.30	1840.75	0	1	1	(
4	0	0	2	1	70.70	151.65	1	0	1	(
...
7038	0	1	24	1	84.80	1990.50	0	1	0	:
7039	0	1	72	1	103.20	7362.90	1	0	0	:
7040	0	1	11	0	29.60	346.45	1	0	0	:
7041	1	1	4	1	74.40	306.60	0	1	1	(
7042	0	0	66	1	105.65	6844.50	0	1	1	(

7043 rows × 43 columns



In [48]: #d=r.get_dummies(d)

```
In [49]: x=r.get_dummies(x)
x
```

Out[49]:

	SeniorCitizen	Partner	tenure	PhoneService	MonthlyCharges	TotalCharges	gender_Female	gender_Male	Dependents_No	Dependents_Yes
0	0	1	1	0	29.85	29.85	1	0	1	(
1	0	0	34	1	56.95	1889.50	0	1	1	(
2	0	0	2	1	53.85	108.15	0	1	1	(
3	0	0	45	0	42.30	1840.75	0	1	1	(
4	0	0	2	1	70.70	151.65	1	0	1	(
...
7038	0	1	24	1	84.80	1990.50	0	1	0	:
7039	0	1	72	1	103.20	7362.90	1	0	0	:
7040	0	1	11	0	29.60	346.45	1	0	0	:
7041	1	1	4	1	74.40	306.60	0	1	1	(
7042	0	0	66	1	105.65	6844.50	0	1	1	(

7043 rows × 43 columns

```
In [66]: from sklearn.model_selection import train_test_split
x_train,x_test,y_train,y_test=train_test_split(x,y,test_size=0.33,random_state=42)
```

```
In [51]: from sklearn.model_selection import GridSearchCV #GridSearchCV is for parameter tuning
from sklearn.ensemble import RandomForestClassifier
cls=RandomForestClassifier()
n_estimators=[25,50,75,100,125,150,175,200] #number of decision trees in the forest, default = 100
criterion=['gini','entropy'] #criteria for choosing nodes default = 'gini'
max_depth=[3,5,10] #maximum number of nodes in a tree default = None (it will go till all possible nodes)
parameters={'n_estimators': n_estimators, 'criterion':criterion, 'max_depth':max_depth} #this will undergo 8*2
RFC_cls = GridSearchCV(cls, parameters)
RFC_cls.fit(x_train,y_train)
```

```
Out[51]: GridSearchCV(estimator=RandomForestClassifier(),
                      param_grid={'criterion': ['gini', 'entropy'],
                                   'max_depth': [3, 5, 10],
                                   'n_estimators': [25, 50, 75, 100, 125, 150, 175, 200]})
```

In a Jupyter environment, please rerun this cell to show the HTML representation or trust the notebook.

On GitHub, the HTML representation is unable to render, please try loading this page with nbviewer.org.

```
In [56]: RFC_cls.best_params_
```

```
Out[56]: {'criterion': 'entropy', 'max_depth': 10, 'n_estimators': 200}
```

```
In [57]: cls=RandomForestClassifier(n_estimators=200,criterion="entropy",max_depth=10)
```

```
In [58]: cls.fit(x_train,y_train)
```

```
Out[58]: RandomForestClassifier(criterion='entropy', max_depth=10, n_estimators=200)
```

In a Jupyter environment, please rerun this cell to show the HTML representation or trust the notebook.

On GitHub, the HTML representation is unable to render, please try loading this page with nbviewer.org.

```
In [61]: rfy_pred=cls.predict(x_test)
```

```
In [62]: rfy_pred
```

```
Out[62]: array([1, 0, 0, ..., 1, 0, 0])
```

```
In [64]: from sklearn.metrics import confusion_matrix  
confusion_matrix(y_test, rfy_pred)
```

```
Out[64]: array([[1548, 149],  
               [ 293, 335]])
```

```
In [65]: from sklearn.metrics import accuracy_score  
accuracy_score(rfy_pred, y_test)
```

```
Out[65]: 0.8098924731182796
```

```
In [68]: from sklearn.linear_model import LogisticRegression  
classifier = LogisticRegression()  
classifier.fit(x_train, y_train)
```

```
Out[68]: LogisticRegression()  
In a Jupyter environment, please rerun this cell to show the HTML representation or trust the notebook.  
On GitHub, the HTML representation is unable to render, please try loading this page with nbviewer.org.
```

```
In [69]: y_pred = classifier.predict(x_test)  
y_pred
```

```
Out[69]: array([1, 0, 0, ..., 1, 0, 0])
```

```
In [70]: from sklearn.metrics import confusion_matrix  
confusion_matrix(y_test, y_pred)
```

```
Out[70]: array([[1538, 159],  
               [ 277, 351]])
```

```
In [71]: from sklearn.metrics import accuracy_score  
accuracy_score(y_pred,y_test)
```

```
Out[71]: 0.8124731182795699
```

```
In [ ]:
```